

AREAS OF INTEREST	AI Alignment (Machine Learning Transparency, Value Learning), Theory of Artificial Intelligence (Reinforcement Learning, Algorithmic Information Theory, Statistical Machine Learning), Economics (Agency Theory)	
DEGREES	<p><i>Doctor of Philosophy in Computer Science</i>, 2016 – present University of California, Berkeley</p> <ul style="list-style-type: none"> Studying AI alignment, supervised by Stuart Russell. Researcher at the Center for Human-Compatible AI. GPA: 3.35/4.00 <p><i>Bachelor of Philosophy (Hons)</i>, 2012 – 2015 Australian National University</p> <ul style="list-style-type: none"> Honours in Computer Science, undergraduate studies in Mathematics and Physics. Thesis: “Resource-bounded Complexity-based Priors for Agents”, supervised by Marcus Hutter. GPA: 7.00/7.00, 1st Class Honours. 	
PUBLICATIONS	<ul style="list-style-type: none"> Loss Bounds and Time Complexity for Speed Priors. With Jan Leike and Marcus Hutter. AISTATS 2016. Self-modification of Policy and Utility Function in Rational Agents. With Tom Everitt (lead author), Mayank Daswani, and Marcus Hutter. AGI 2016, recipient of Kurzweil Prize for Best Paper. Exploring Hierarchy-Aware Inverse Reinforcement Learning. With Chris Cundy (lead author). GoalsRL Workshop at ICML/IJCAI/AAMAS 2018. 	
SELECTED AWARDS	<p><i>University Medal</i>, Australian National University 2015</p> <ul style="list-style-type: none"> Prize; awarded to students who have obtained First Class Honours (or Masters Advanced Equivalent) and demonstrated exceptional academic excellence across their studies, the highest academic prize for undergraduates. <p><i>Erin Brent Computer Science Prize</i>, Australian National University 2015</p> <ul style="list-style-type: none"> Monetary prize; awarded to the student who achieved the best Honours result in any of the degree programs relating to Computer Science, Software Engineering or Information Technology. 	
INTERNSHIPS	<p><i>Machine Intelligence Research Internship</i> 2019</p> <ul style="list-style-type: none"> Spent 3 months on research engineering team 4 days per week, while supervising a UC Berkeley intern 1 day a week. <p><i>Future of Humanity Institute, Oxford University</i> 2016</p> <ul style="list-style-type: none"> Writing code for agentmodels.org, a website designed to explain the use of probabilistic programs to build models of agents and perform inference about them. 	